

# CSE 398/498 Lab 1 Report:

## Object Detection with YOLOv11 and YOLOE

Hengde Dai  
hed424@lehigh.edu

September 2, 2026

### 1 Overview

I completed the three experiments of Lab 1: (1) ran pretrained YOLOv11 on my laptop webcam, (2) fine-tuned YOLOv11 on a 50-image custom dataset and redid the webcam test, and (3) followed the YOLOE Colab to learn open-set detection. All code, data, weights, and executed notebooks are in my GitHub repository.<sup>1</sup>

Environment: webcam experiments ran locally (Windows 11, RTX 3070 Ti, PyTorch 2.5.1 + CUDA 12.1); training and the YOLOE notebook ran on a Colab T4 GPU.

### 2 Experiment 1: Pretrained YOLOv11 on a Webcam


I cloned the assigned `webcam-object-recognition` repo and tested nine household objects in front of the camera. The script uses COCO-pretrained `yolo11n.pt` with a 0.5 confidence threshold. Table 1 lists the results; Figure 1 shows one success and the failure that I later used for fine-tuning.

| Object                 | Prediction | Confidence | Outcome                                 |
|------------------------|------------|------------|-----------------------------------------|
| Water bottle           | bottle     | 0.91       | correct                                 |
| Couch                  | couch      | 0.85       | correct                                 |
| Scissors               | scissors   | 0.80       | correct                                 |
| Book                   | book       | 0.78       | correct                                 |
| Cell phone             | cell phone | 0.78       | correct                                 |
| Game controller        | remote     | 0.82       | near miss (no controller class in COCO) |
| Vine tomatoes          | apple      | 0.52–0.62  | wrong                                   |
| Toilet paper (upright) | cup        | 0.59       | wrong                                   |
| Toilet paper (on side) | toilet     | 0.53       | wrong                                   |


Table 1: Pretrained YOLOv11n webcam results. Screenshots are in `step1-webcam-yolov11/results-pretrained/`.

The pattern behind the failures is simple: COCO has 80 classes, and tomato and toilet paper are not among them. The model can only pick the closest class it knows, so tomatoes became “apple” and a paper roll became “cup” or “toilet.”

<sup>1</sup><https://github.com/DylanDai63/cse498-lab1>



(a) Correct: bottle, 0.91.



(b) Wrong: tomatoes detected as apple.

Figure 1: Pretrained YOLOv11n on the webcam.

### 3 How YOLOv11 Works

YOLO means “You Only Look Once”: one CNN forward pass over the whole image directly outputs boxes, confidence scores, and class scores. Older two-stage detectors (e.g., R-CNN) first propose regions and then classify them; YOLO’s single-stage design is less accurate but much faster, which is why the webcam demo runs in real time (about 7 ms per frame on my GPU).

Key ideas:

- **Grid prediction.** The image is divided into an  $S \times S$  grid. Each cell predicts a few boxes  $(x, y, w, h)$  with confidence plus class scores. The cell containing an object’s center is responsible for that object.
- **Training assignment.** Each ground-truth object is matched to the predicted box with the highest IoU, so one predictor specializes per object.
- **NMS.** At inference many overlapping boxes appear. Non-maximum suppression keeps the highest-confidence box and drops neighbors whose IoU with it exceeds a threshold.
- **Architecture.** Backbone (feature extraction) + neck (multi-scale feature fusion) + head (box/class outputs).

YOLOv11 improvements over v8 (from the course slides): the C3k2 block replaces C2f in the backbone for cheaper multi-scale features; the C2PSA attention module helps the network focus on informative regions; the head shares weights between classification and regression branches to cut latency. Net effect: higher COCO accuracy with fewer parameters. I used the nano variant yolo11n (2.6M parameters) in all experiments so the comparison stays fair.

## 4 Experiment 2: Fine-Tuning YOLOv11

### 4.1 Dataset (50 images, class tomato)

The Kaggle grocery dataset suggested in the handout turned out to be a CSV of product listings with no images, so, following the revised instructions, I built my own dataset from Google Open Images. I picked `tomato` as the target and checked programmatically that it is not among the 80 COCO class names, so it is an unseen class. The final dataset has 50 images / 99 boxes in YOLO format, split 42 train / 8 validation (Figure 2).

It took me three tries to get a dataset that works:

1. **Two classes (tomato + toilet paper):** tomato reached mAP50 0.59 but toilet paper only 0.01. Open Images has very few toilet-paper images; about 12 training images were not enough to learn a class.
2. **50 random tomato images:** mAP50 0.66, but every detection scored below 0.25 confidence. The random sample was noisy: Open Images labels sauces, slices, unripe fruit, and even a persimmon photo as “Tomato.”
3. **50 hand-picked images:** I ranked candidates by box size, reviewed about 200 by hand, and kept 50 clean photos of whole, mostly red tomatoes. I also excluded dense piles where not every visible tomato is labeled, because unlabeled objects count as background during training.




Figure 2: A training batch with Ultralytics augmentation (mosaic, flips, color jitter).

## 4.2 Training and Results

I followed the assigned Roboflow Colab, with three changes: model `yolo11n.pt` (to match Experiment 1), my own `data.yaml`, and `epochs=150` (a 42-image set needs many passes; an earlier 50-epoch run gave correct boxes but with confidence too low to pass normal thresholds). Training took about 5 minutes on the T4.

Results on the 8-image validation set: precision 0.788, recall 0.714, mAP50 0.727, mAP50-95 0.607. The validation set is small, so these numbers are rough indicators. Figure 3 shows the training curves; Figure 4 shows the confusion matrix (12 of 21 tomatoes detected, 9 missed, 3 false alarms) and sample predictions.

## 4.3 Redoing Experiment 1: Does Fine-Tuning Help?

Yes. I loaded `best.pt` in the same webcam script and held up the same vine tomatoes: the pretrained model had called them `apple` (0.62); the fine-tuned model says `tomato` at 0.96 (Figure 5). One side effect: the fine-tuned model detects nothing else. It only knows the one class it was trained on and forgot the 80 COCO classes.






Figure 3: Losses (left) and validation metrics (right) over 150 epochs.



(a) Confusion matrix (0.25 threshold).



(b) Predictions on validation images.

Figure 4: Fine-tuned model on the validation set.

#### 4.4 What I Learned from Fine-Tuning

- **Data quality beats quantity.** 50 noisy images gave a model that was useless in practice; 50 curated images fixed it with the same training recipe.
- **Each class needs a minimum amount of data.** Toilet paper failed with about 12 training images; tomato worked with 42. Transfer learning lowers the requirement but does not remove it.
- **Check metrics together with a threshold.** My intermediate model had decent mAP (which averages over all thresholds) yet detected nothing at 0.25. Only the confusion matrix revealed the problem.
- **Fine-tuning on new classes forgets old ones.** To keep both, one would include old classes in the training data or keep two models.




Figure 5: After fine-tuning: the same tomatoes, now `tomato 0.96`.

## 5 Experiment 3: Open-Set Detection with YOLOE

### 5.1 What Is Open-Set Object Detection?

A normal (closed-set) detector can only output the classes fixed at training time; my two models show the problem from both sides (80 classes that miss tomato, or one class that misses everything else). Open-set (open-vocabulary) detection removes the fixed list: you give the model a *prompt* at inference time, as text (class names) or as a visual example (a reference box), and it finds matching objects even for categories it was never explicitly trained on.

In the assigned YOLOE Colab, the text prompt `["dog", "eye", "nose", "tongue"]` made the model detect all four concepts on the example image, including object parts (Figure 6). No retraining was involved; changing what the model detects only takes editing the word list.




Figure 6: YOLOE text-prompt detection, zero-shot.

### 5.2 YOLOE Architecture and What Changed

YOLOE keeps the standard YOLO skeleton: a YOLOv8/v11-style backbone and a PAN neck that fuses the P3/P4/P5 scales, so it stays real-time. The key change is in the head:

- A closed-set YOLO head ends in a classification layer with one fixed weight row per training class.

- YOLOE replaces it with an anchor-free, decoupled head whose classification branch outputs an **embedding**, and each region is scored by **similarity against prompt embeddings**. The class list becomes an input instead of a trained parameter.
- Three prompt encoders produce the embeddings: **RepRTA** for text (class names go through a CLIP-style text encoder once, and the resulting vectors act as the classification weights), **SAVPE** for visual prompts (a reference box is encoded into the same space), and **LRPC** for prompt-free mode (a built-in vocabulary of about 4,585 classes).

Why it works: during large-scale pretraining on region-text data, the detector learns to place each region’s feature close to the text embedding of its label. A new word from the user lands at a meaningful spot in the same space, so similarity matching still works. A plain COCO-trained YOLO cannot do the trick, because its features were never aligned with a text space; the alignment pretraining is exactly what YOLOE adds. Once you fix a word list, the embeddings can be folded back into a normal classification layer (re-parameterization), so deployment costs nothing extra.

Compared with Experiment 2: teaching my closed-set model one new word took a curated dataset and 150 epochs; YOLOE handles new words from one line of text, at the price of a bigger model and an expensive pretraining paid once by its authors.

## AI Use Statement

Per the course AI policy: I used Claude (Anthropic) as an assistant for environment setup and debugging, building and curating the Open Images dataset, step-by-step guidance through the Colab notebooks, and drafting and revising the report text. I ran all experiments myself and take responsibility for all content, code, and results.